

Assignment 2

(Dated: May 31, 2021)

In this problem set, we will begin to explore **Bilby**, which is a package that provides a consistent framework to run different types of MCMC and nested samplers on generic distributions. Bilby itself is not a sampler, but wraps popular samplers like **emcee** and **dynesty** with an additional layer to make the analysis as modular as possible. Bilby is not specific to GW science, but can be used to sample from any probability distribution with any of the supported samplers. To begin with, you will explore three tutorials (2 from Bilby website and 1 from Deep). You should download both examples, run them, and be familiar enough with the code that you could modify the likelihood functions and samplers.

1. The first **tutorial** fits a 1-D gaussian curve to a series of data drawn from a normal distribution. In this context, the two parameters of the model are the mean (μ) and standard deviation (σ) of the distribution. That is, our model is defined by the likelihood function

$$p(x|\mu, \sigma) = \frac{1}{\sqrt{\sigma^2 2\pi}} e^{-\frac{1}{2} \frac{(x-\mu)^2}{\sigma^2}}, \quad (1)$$

where x is our data. The analysis begins by defining a class for the likelihood of interest. This format is the standard way Bilby works: you define a class object with the specific members (like the `log_likelihood` function in this tutorial), and that provides enough abstraction for Bilby to pass this likelihood on to the sampler of choice. Bilby also allows you to set the priors for the parameters as well, which, for this tutorial, are uniform from $[0, 5]$ and $[0, 10]$ for μ and σ , respectively. The priors used in this tutorial are pre-made in Bilby, and are in the `bilby.core.prior` file, along with other common choices of priors. The priors can also be defined in a similar way as the likelihood, but custom priors will be addressed later. With the problem defined by the full posterior (prior + likelihood), the sampler can be run with `bilby.run_sampler` and the results are plotted in a corner plot by `result.plot_corner()`.

First, read into the documentation of Bilby and **dynesty** to understand what the arguments of `run_sampler` do. Run the sampler for multiple different configurations and examine the output plot. Some additional modifications you can explore is to draw samples from a non-gaussian distribution for the data (say student-t) and fit a gaussian to it, or modify the priors such that they are not uniform.

2. The second **tutorial** walks you through performing linear regression on a set of data with noise of unknown variance. In this tutorial, the model assumes the data is linear (with 2 parameters, the slope and the intercept), and that the noise is gaussian, but with unknown variance (with 1 additional parameter, on top of the signal model). The likelihood function in this tutorial uses the predefined Gaussian likelihood class already included in Bilby, so you should read into that class and understand the arguments to `bilby.core.likelihood.GaussianLikelihood`. The priors are again uniform for the 3 parameters. Again, run the tutorial and understand each line of it. Then, modify the parameters, rerun the analysis, and examine the input.
3. Run Deep's custom-prior tutorial